TenderMaster

PRODUCT BLUEPRINT & MARKET RESEARCH · V1.0

The AI that reads, qualifies & compiles winning tenders.

An MVP blueprint for an AI automation SaaS that turns 200-page EPC tenders into a clear eligibility verdict, a live document checklist, and a single submission-ready PDF — built around a company knowledge graph that is genuinely hard to copy.

EPC / Construction · Pakistan

Single plan · $150 / month

Eligibility-first AI

ERP + Reminder layer

CONTENTS

What's inside

1 · How EPC tendering actually works in Pakistan

Most EPC (Engineering, Procurement & Construction) work flows through **public-sector procurement**, regulated by **PPRA** — federally (PP Rules 2004) and through provincial bodies (PPRA Punjab, SPPRA Sindh, KPPRA, BPPRA). Contracts are floated as EPC / Turnkey / LSTK packages, almost always under **FIDIC** conditions (Silver Book for turnkey). The default bidding method is **Single-Stage Two-Envelope (SSTE)**; large/complex jobs use Two-Stage Two-Envelope.

The lifecycle a bidder lives through

1. **Discovery.** Tender advertised on EPADS / PPRA portals, in newspapers, and on the agency's own site. Bidder screens for fit.
2. **Eligibility check.** Does our PEC category, turnover, and similar-project experience meet the bar? This is the highest-leverage decision — and where the product starts.
3. **Document purchase.** Buy/download the bidding document (a fee, often Rs 2,000–25,000).
4. **Bid security.** Arrange a bank guarantee / CDR / pay-order, typically **1–5%** of bid value.
5. **Compile.** Assemble Technical Bid (no prices) + Financial Bid (priced BOQ) into two sealed envelopes / two e-uploads.
6. **Submit before deadline.** Via EPADS (federal, mandatory above threshold) or physically. Submission day = opening day.
7. **Public opening & evaluation.** Technical envelopes opened first; only technically-responsive bidders' financial envelopes are opened. Award to lowest-evaluated responsive bid.
8. **Award & mobilise.** Winner posts a Performance Guarantee (5–10%) and signs the contract.

> **The product's wedge:** 80% of a bidder's pain lives in steps 2 & 5 — "Am I even eligible?" and "Did I attach every single required document, correctly?". A missing affidavit or an out-of-category PEC licence gets a bid rejected before price is ever read. That is the problem TenderMaster exists to kill.

2 · Eligibility, registrations & the PEC category ladder

Before a firm can legally bid, several credentials must exist and be current. The AI must **store, verify, and watch the expiry of every one of these** per company — they are the backbone of every eligibility check.

CREDENTIAL	ISSUED BY	WHY IT GATES THE BID
PEC Constructor Licence	Pakistan Engineering Council	Mandatory to execute engineering works. The tender names a **minimum category**; below it = auto-reject.
SECP / Firm registration	SECP / trade office	Legal entity proof.
NTN + Sales-Tax reg. + Active Filer	FBR	Tax compliance; non-filers are excluded.
EPADS vendor account	PPRA	Required to e-submit federal bids.
Bank standing / line of credit	Bank	Backs bid security & proves liquidity.
Sector licences (NEPRA/AEDB, PSQCA, ISO)	Various	Power, water, quality-specific gates.

The PEC category ladder — the single biggest gate

Your licence category **caps the project value you may bid for**. Firms climb it over years by accumulating completed-project value and PCP (Professional Credit Points). The AI must map every tender's value to the required category and instantly tell the firm if they're under-categorised.

CATEGORY	INDICATIVE PROJECT-VALUE CEILING	TYPICAL FIRM
C-6	~Rs 25 million (entry; non-engineers allowed)	New / micro contractor
C-5 → C-3	Rising ceilings (tens → hundreds of millions)	Growing SME EPC firm
C-2 / C-1	Large (hundreds of millions+)	Established contractor
C-B / C-A	Very large / **no limit**	Top-tier national EPC

> **Data point worth owning:** the exact, current category ceilings change with PEC policy. TenderMaster should treat these as versioned reference data it maintains centrally — every customer benefits the instant it's updated. (See Moat, Layer 4.)

What "eligibility" really means for one tender

- **Hard gates** (pass/fail): PEC category ≥ required · active filer · valid registrations · turnover ≥ X · similar-work experience ≥ N projects of ≥ Rs Y.
- **Soft / scored**: equipment owned, key personnel quality, past performance, financial ratios, completion certificates.

3 · Anatomy of a tender & the master document checklist

Every bidding document, however long, decomposes into the same skeleton. The AI's **extraction job** is to read the PDF and pull these into structured fields, then derive the checklist.

What's in the tender document

- Invitation for Bids (IFB) — scope, deadline, fee
- Instructions to Bidders (ITB)
- **Eligibility & qualification criteria**
- Scope of work / technical specs / drawings
- Bill of Quantities (BOQ) / price schedule
- Conditions of Contract (FIDIC-based)
- Bid & performance security clauses
- **Evaluation & scoring criteria**
- Forms, affidavits, integrity pact

The two-envelope bid we must build

Envelope A — Technical (no prices)

Envelope B — Financial (priced BOQ + bid security)

Technical is opened & scored first; only the responsive bidders' financial envelopes get opened. The compiled PDF TenderMaster generates **is Envelope A** — the document-heavy half where bids die.

Master document checklist (the AI's reusable backbone)

Every tender selects a subset of these. The AI maps tender language → checklist items → the company's stored documents, then flags gaps.

- Covering letter / bid form (signed, stamped)
- PEC constructor licence (correct category)
- SECP / firm registration certificate
- NTN & Sales-Tax registration; Active-filer printout
- Audited financial statements (3 yrs)
- Bank statement / line-of-credit letter
- Turnover certificate
- Bid security (BG / CDR / pay-order)
- List of similar completed projects
- Completion certificates / client PCs
- Work-in-hand / commitments statement
- Key personnel CVs + PEC reg. of engineers
- Owned / leased plant & equipment schedule
- Methodology & work programme (bar chart)
- Affidavit — non-blacklisting
- Integrity pact
- Power of attorney / authorised signatory
- JV agreement (if applicable)
- Compliance / specification response sheet
- ISO / sector certifications

Evaluation: how the bid is scored

Technical responsiveness is **pass/fail against the hard gates**; many tenders then apply a **weighted technical score** (experience, personnel, methodology, equipment, financial strength). Lowest-evaluated responsive price wins. **Knowing the weightings lets the AI tell a firm not just "you qualify" but "you'd score ~72/100 — here's the 18 points you're leaving on the table."**

4 · The 7 real pain points we remove

These are the concrete jobs EPC firms hate. Each maps to a TenderMaster module — design the product around the pain, not the feature.

#	THE PAIN TODAY	WHAT IT COSTS THEM	TENDERMASTER REMOVES IT VIA
1	Reading a 150–300 page tender to find the 12 things that actually matter	2–6 hours of a senior engineer, per tender	AI tender reader → structured summary
2	Deciding "are we even eligible?" — often discovered after effort is spent	Wasted bids, lost bid-doc fees	Eligibility / Go–No-Go engine
3	Figuring out the exact document list and chasing each one	Rejected on a missing affidavit	Dynamic checklist + upload + "acquire" help
4	Re-typing the same company data into every bid form	Hours of copy-paste, transcription errors	Knowledge Graph auto-fill
5	Compiling 20+ PDFs into one correctly-ordered submission	Last-night panic, wrong order = penalised	One-click compiled PDF generator
6	Tracking deadlines, securities expiry, validity windows across many live tenders	Missed deadlines = zero recovery	ERP + reminder engine
7	No memory of what they bid, won, or lost — and why	Same mistakes repeat; no learning	Tender history + win/loss feedback loop

150–300
PAGES PER TENDER

20+
DOCUMENTS TO COMPILE

1 missing doc
= FULL REJECTION

90–120d
BID VALIDITY TO TRACK

> **Reframing the value:** the firm isn't paying $150/mo for "AI that fills forms." They're paying to never lose a winnable tender to a clerical mistake again, and to bid on more tenders with the same staff. That's the pitch.

5 · Vision, principle & the six core modules

Operating principle (your instinct, sharpened): TenderMaster is not a chat assistant. It is a **verdict-and-checklist engine**. It reads a tender and tells the firm three things in order: **(1) Can you bid? (2) What do you need? (3) Here is your finished package.** Conversation is a tool inside that flow, never the product.

The six core modules

MODULE 1

Company Knowledge Graph

An extensive onboarding builds a living, structured profile of the firm — registrations, financials, projects, people, equipment, documents. Everything downstream reads from this. This is the asset.

MODULE 2

Tender Reader & Extractor

Upload/import a tender PDF → AI returns a structured brief: scope, deadline, required category, securities, evaluation weights, and every requirement as a discrete item.

MODULE 3

Eligibility / Go–No-Go Engine

Cross-checks tender hard-gates against the Knowledge Graph → a clear verdict: **Eligible / Eligible-with-gaps / Not eligible**, with the exact reasons and a predicted technical score.

MODULE 4

Dynamic Document Checklist

Per-tender checklist; each row has status (have / missing / expired), an **upload** slot, auto-fill from the graph, and an **"Ask AI to help me acquire this"** action.

MODULE 5

Compiled PDF Generator

One click → a correctly-ordered, indexed, bookmarked submission PDF (Envelope A), plus an optional **"edge" section** of extra strengths not required but likely to win marks.

MODULE 6

ERP · Reminders · History

Tracks every live tender, deadline, security expiry; nudges the user ("did you submit?", "deadline in 48h", "licence expires in 30d"); archives every past tender with outcome.

How they connect

- → **Knowledge Graph** feeds **Eligibility**, **Checklist auto-fill**, and **PDF generation**.
- → **Tender Reader** output drives the **Eligibility verdict** and the **Checklist contents**.
- → **ERP** wraps everything in time — deadlines, reminders — and the **History/feedback loop** teaches the graph what wins.

6 · The Company Knowledge Graph (onboarding)

This is the heart of the moat. The onboarding is **deliberately thorough** — a one-time, guided, multi-session intake that produces a structured profile rich enough that the AI never has to ask the firm a basic question again. Below is the full data model the MVP should capture.

1 · Legal & Identity

- Registered name, brand name, type (Pvt Ltd / sole / partnership / JV)
- SECP reg #, incorporation date
- NTN, STRN, filer status
- Registered & site addresses, contacts
- Authorised signatories + power of attorney

2 · Licences & Registrations

- PEC category, licence #, **validity / expiry**
- Agency enlistments (NHA, WAPDA, C&W, cantonments…)
- Sector regs (NEPRA, AEDB, PSQCA), ISO certs
- Chamber of commerce membership

3 · Financial Profile

- Annual turnover (last 3–5 yrs)
- Audited statements, auditor name
- Bank(s), line of credit, available limit
- Working capital, net worth, key ratios
- Bid-security capacity

4 · Experience / Projects

- Per project: name, client, value, scope, sector
- Start/end dates, status, role (prime/sub/JV)
- Completion certificate (uploaded)
- Tags: discipline, capacity (MW, km, m^3…)
- Work-in-hand / current commitments

5 · People & Capability

- Key personnel: name, role, qualification, PEC reg, years exp.
- CVs (uploaded, parsed)
- Org chart / total headcount by discipline
- In-house disciplines vs outsourced

6 · Plant, Equipment & Documents

- Owned/leased equipment schedule (type, qty, capacity)
- Document vault: every certificate, affidavit template, signed form
- Reusable boilerplate: methodology, HSE policy, QA/QC plan
- Logos, letterhead, stamps for generated PDFs

> **Make onboarding feel like progress, not a tax.** Show a live "Tender Readiness Score" that climbs as the firm fills the graph — e.g. **"You're 78% ready. Add 2 completion certificates to unlock bids up to Rs 400M."** Every field they add visibly increases the tenders they can win. That turns a chore into a game and silently builds your moat.

Design note: store everything as typed, queryable entities (not free text) — a project is a row with value/sector/dates, not a paragraph. That is what lets the Eligibility Engine do real matching and what a copycat can't reproduce without re-collecting all the data.

6 · (cont.) Turning the graph into instant answers

Once the graph exists, the same data answers many questions automatically. One profile, many derived outputs:

- **Eligibility match** → tender hard-gates vs graph → pass/fail per gate
- **Auto-filled forms** → bid form, qualification forms, personnel & equipment schedules pull straight from entities
- **Predicted score** → graph mapped onto the tender's evaluation weights → "≈72/100, here's how to reach 84"
- **Gap list** → exactly which documents/credentials are missing or expired for this tender
- **The winning "edge" pack** → strengths in the graph that aren't required but score marks, surfaced automatically
- **Capacity ceiling** → "with your current category & turnover you can safely bid up to Rs X"

The "Ask AI to help me acquire this" action

For any missing checklist item, a contextual helper that does specific, useful things — not generic chat:

It explains

"This tender needs a non-blacklisting affidavit on Rs 100 stamp paper. Here's the exact wording and who must sign it."

It generates

Drafts the affidavit / integrity pact / covering letter / methodology pre-filled from the graph, ready to print & sign.

It routes

"A bank guarantee for bid security is issued by your bank against margin. Steps: 1… 2… Typical time: 2–3 days — start now, deadline is in 9 days."

It reuses

"You already uploaded this completion certificate for Tender #14 — attach the same one?"

7 · End-to-end user journey

What a user actually experiences, from cold signup to a submitted bid and beyond.

1. **Guided onboarding.** Firm builds its Knowledge Graph over a few sittings. Readiness Score climbs to, say, 85%.
2. **New tender in.** User uploads a tender PDF (or it's auto-pulled from EPADS/PPRA in a later version). AI parses it in ~1 minute.
3. **Verdict screen.** Big, unambiguous: ✓ Eligible / ▲ Eligible with 3 gaps / ✕ Not eligible (PEC C-4 required, you hold C-5) — plus a one-paragraph plain-language tender summary and predicted score.
4. **Checklist appears.** Auto-populated. Green = already in your vault, amber = expired/needs refresh, red = missing. Each red row has Upload + "Ask AI to acquire."
5. **Fill the gaps.** User uploads or generates the missing pieces; AI validates each (right category? in date? right name?).
6. **Generate package.** One click → ordered, indexed, bookmarked submission PDF + optional edge pack. Preview, then download.
7. **Track & submit.** Tender moves to "Ready." Reminders fire toward the deadline. User marks it **Submitted**.
8. **Close the loop.** Later, AI asks the outcome (Won / Lost / Shortlisted) and why. That feeds history + the learning loop.

> **The "wow" moment to nail in the demo:** upload a real 200-page tender, and within 90 seconds the user sees a clear Eligible / Not-eligible verdict and a complete checklist with half the rows already green from their vault. If that lands, the product sells itself.

8 · Seven layers that make it hard to copy

A UI is copyable in a weekend. **Accumulated, structured, proprietary intelligence is not.** The defensibility comes from stacking these — any one is weak; together they compound.

1 DATA LOCK-IN **The Company Knowledge Graph.** Once a firm has spent hours encoding every project, certificate and ratio, switching means re-doing all of it. The data is the switching cost. The richer the graph, the stickier the user.

2 DOMAIN RULE ENGINE **A codified Pakistani-procurement rulebook.** PPRA/PEC rules, category ceilings, FIDIC quirks, agency-specific forms, affidavit wordings — encoded as a maintained, versioned rules layer. This is hundreds of hours of domain capture a generic "AI RFP tool" doesn't have.

3 AGENCY INTELLIGENCE **Per-organisation tender memory.** The AI learns how each procuring agency (NHA vs WAPDA vs a cantonment board) words requirements, which documents they always demand, how they score. Every tender processed sharpens this. A copycat starts from zero.

4 REFERENCE DATA **Centrally-maintained ground truth.** Live category ceilings, current fee structures, standard form templates, blacklist watchlists — you maintain them once, every customer benefits instantly. Becomes the authoritative source firms trust.

5 WIN/LOSS LOOP **Outcome learning.** Because the ERP captures what was bid and what won/lost, the system learns which profiles & framings win which tenders — and can advise "firms like you that won this agency added X." Proprietary, compounding, impossible to fake on day one.

6 WORKFLOW DEPTH **It owns the whole lifecycle.** Discovery → eligibility → compile → submit → remind → outcome. A clone can copy one screen; replicating the entire integrated loop (plus reminders that actually fire and a PDF engine that gets ordering right) is a real product, not a wrapper.

7 TRUST & RECORD **System of record.** Once it holds a firm's entire tender history, securities calendar and document vault, it becomes the place they operate from. Leaving means losing their institutional memory. That's the deepest moat of all.

> **Strategic warning:** the moat is earned over time, not shipped on day one. The MVP's job is to start the flywheel — get firms to pour data in and run tenders through — so layers 1, 3, and 5 begin compounding immediately. Build the data-capture loops first; the defensibility follows.

PART D · BUILD IT

9 · MVP scope & AI architecture

Ruthless scoping. Ship the wedge — eligibility + checklist + compiled PDF — backed by a real knowledge graph. Everything else waits.

In the MVP V1

- Guided onboarding → Knowledge Graph
- Tender PDF upload + AI extraction
- Eligibility / Go–No-Go verdict + plain summary
- Dynamic checklist (have/missing/expired)
- Per-row upload + "Ask AI to acquire"
- Document generation (affidavits, letters, schedules)
- One-click compiled submission PDF + edge pack
- Tender dashboard + deadline reminders (email)
- Tender history with outcomes
- Single plan, $150/mo, dark-purple theme

Deliberately later V2+

- Auto-discovery: scrape EPADS/PPRA & match to firm
- Agency-intelligence & win/loss recommendation engine
- Predicted-score optimiser ("reach 84/100")
- Multi-user / team roles & approvals
- WhatsApp/SMS reminders
- BOQ pricing assistance (financial envelope)
- Direct e-submission integrations
- Mobile app

AI architecture (how the intelligence is built)

1 · Ingestion & OCR. Tender PDF → text + layout. Many Pakistani tenders are scanned — OCR (and Urdu/English handling) is mandatory, not optional.

2 · Structured extraction. LLM with a strict schema pulls requirements into typed fields (deadline, PEC category, securities, each eligibility gate, evaluation weights). Use function-calling / JSON-schema output — not free-form.

3 · Retrieval (RAG). Index the company graph + your domain rulebook + agency memory; retrieve the relevant slices to ground every verdict and generated document. Grounded, cited, low-hallucination.

4 · Rules + reasoning split. Hard gates run through deterministic rules (category ≥ X is code, not vibes); only soft judgement & drafting use the LLM. This keeps eligibility verdicts trustworthy.

5 · Generation. LLM drafts documents from graph data into your templates; a deterministic engine assembles the final ordered PDF.

6 · Feedback capture. Outcomes & corrections logged → improves extraction prompts, rules, and agency memory over time.

> **Model choice:** default to the latest, most capable Claude models (e.g. Claude Opus 4.8 for hard extraction/reasoning, a faster Claude for routine drafting) behind a provider-agnostic layer. Reliable structured extraction and grounded generation matter more here than raw speed — a wrong eligibility verdict

10 · Data model & suggested stack

Core entities (the schema behind the graph)

ENTITY	KEY FIELDS	FEEDS
Company	legal id, type, addresses, NTN/STRN, filer status	everything
Credential	type (PEC/ISO/…), number, category, issue/expiry, file	eligibility, reminders
Project	client, value, sector, role, dates, capacity, completion cert	experience gates, score
Person	role, qualification, PEC reg, years, CV	personnel scoring
Equipment	type, qty, owned/leased, capacity	equipment schedule
Financial	year, turnover, net worth, ratios, audited file	financial gates
Document	type, file, version, linked entity	vault, checklist fill
Tender	agency, value, deadline, required category, extracted requirements (JSON), evaluation weights	eligibility, checklist
Submission	tender ref, checklist state, generated PDF, status, outcome, reason	history, learning loop
Requirement	tender ref, text, type (gate/doc/scored), matched-to, status	checklist rows

Suggested stack (lean, MVP-appropriate)

Frontend

Next.js / React · Tailwind (dark-purple design system) · component lib for forms-heavy onboarding

Backend

Node/TypeScript or Python (FastAPI) · Postgres (relational graph + JSONB for extracted requirements) · object storage for files

AI layer

Claude API for extraction & generation · vector store (pgvector) for RAG · OCR (Textract / Tesseract / Google Vision) · deterministic rules service

PDF engine

Headless Chromium / Puppeteer for HTML→PDF · pdf-lib for merge, order, bookmarks, page numbers, index

Jobs & reminders

Queue (BullMQ / Celery) · scheduler for deadline & expiry reminders · email (Resend/SES); WhatsApp later

Infra & auth

Single-tenant-per-company isolation · auth + RBAC-ready · audit log · encrypted document vault

> **Non-negotiable:** the vault holds a firm's most sensitive papers. Encryption at rest, strict per-company isolation, and an audit trail are table stakes for trust — and a selling point against a fly-by-night clone.

11 · My recommendations to make it win more

You asked for suggestions. Here are the ones I'd fight for — ordered by impact.

① Lead with the verdict, not the chat.

Your instinct to avoid "personal assistant" framing is right and is your differentiator. Make the **Eligibility verdict the hero screen**. Every competitor is a chatbot; you're a decision engine. Sell certainty.

② Make onboarding a scored, staged game.

A live Tender Readiness Score + "unlock bids up to Rs X" turns the tedious data capture into visible progress — and it's exactly what builds the moat. Don't bury it; celebrate it.

③ Ship a "predicted score + gap-to-win."

Going beyond pass/fail to **"you'd score ~72/100; add these 3 things to reach 84"** is a feature no generic tool offers and is hard to copy because it needs the evaluation-weight intelligence. High-value, defensible.

④ Capture outcomes religiously.

The win/loss loop is your compounding advantage. Make marking outcomes frictionless (one tap from a reminder). Every outcome makes the next firm's advice better — that's the flywheel.

⑤ The "edge pack" is your emotional hook.

Auto-surfacing extra, non-required strengths that score marks is the part that makes a firm feel the AI is fighting for them. Lean into it — it's also a natural upsell story later.

⑥ Start single-sector, go deep, then widen.

Nail **EPC/construction in Pakistan** before touching other sectors or countries. Depth is the moat; breadth dilutes it early. The same engine extends to IT, supplies, services tenders later — but only after you own one vertical.

⑦ Treat reminders as a product, not a feature.

Missed deadlines are the #1 silent killer. Aggressive, smart, multi-channel nudges (deadline, security expiry, validity lapse, licence renewal) alone can justify the subscription. This is your retention engine.

> **One more:** add a lightweight "Tender Health" badge per live bid (e.g. 6/9 documents ready, 4 days left, on track / at risk). Founders of EPC firms manage by glance — give them a dashboard they check every morning, and you become indispensable.

12 · Pricing logic, roadmap & risks

The single $150/month plan

One plan keeps positioning sharp and sales simple. To make $150/mo feel obviously worth it, anchor it against the cost of one lost winnable tender (often crores of rupees) and the senior-engineer hours saved per bid. Watch your **per-tender AI/OCR cost** — heavy tenders are token-expensive; a soft fair-use cap (e.g. N tenders/month, generous) protects margin without adding plan complexity.

$150	1 plan	ROI
PER MONTH, ONE PLAN	NO TIER CONFUSION	= 1 SAVED BID

Phased roadmap

PHASE	FOCUS	OUTCOME
0 · Foundation	Knowledge Graph + onboarding + document vault	The asset exists; data starts flowing in
1 · Wedge (MVP)	Tender reader + eligibility + checklist + compiled PDF + reminders	Sellable product; flywheel starts
2 · Intelligence	Predicted score, agency memory, win/loss loop	Moat deepens; advice gets unique
3 · Reach	Auto-discovery, team mode, more sectors	Scale & expansion

Top risks & how to manage them

RISK	MITIGATION
Wrong eligibility verdict destroys trust	Deterministic rules for hard gates; always show why; "verify-before-submit" disclaimer; human-in-loop on edge cases
Scanned/messy tender PDFs break extraction	Strong OCR; confidence flags; let users confirm/correct extracted fields (which also feeds learning)
Per-tender AI cost erodes margin	Cache, chunk smartly, route easy steps to cheaper models, fair-use cap
Slow onboarding → churn before value	Readiness Score, templates, import helpers, "minimum viable graph" to run a first tender fast
Data sensitivity / breach	Encryption, isolation, audit log, clear data policy — turn security into a selling point
Copycats	Race to accumulate graph + agency memory + outcomes; depth in one vertical; own the system-of-record relationship

A · The onboarding questionnaire

A · The onboarding questionnaire (MVP intake)

Section 1 — Identity & legal

- Registered & brand name, entity type
- SECP reg #, incorporation date
- NTN, STRN, current filer status
- Head office & site addresses, key contacts
- Authorised signatory + power of attorney

Section 2 — Licences

- PEC category + licence # + expiry (upload)
- Agency enlistments (which, level)
- Sector regs / ISO (upload each)
- Chamber membership

Section 3 — Financials

- Turnover, last 3–5 yrs (upload audited a/cs)
- Bank(s) + line of credit + available limit
- Net worth, working capital
- Max bid-security you can currently post

Section 4 — Projects (repeatable)

- Name, client, value, sector, role
- Start/end, status, capacity metric
- Completion certificate (upload)
- Current work-in-hand

Section 5 — People & plant

- Key personnel: role, qualification, PEC reg, years (upload CV)
- Headcount by discipline
- Owned/leased equipment schedule

Section 6 — Vault & branding

- Standard affidavits, integrity-pact templates
- Methodology / HSE / QA-QC boilerplate
- Logo, letterhead, stamp, signature image

APPENDIX B

B · Eligibility scorecard & closing note

B · Sample eligibility scorecard (what the engine outputs)

● Pass ● Gap — fixable ● Fail — hard stop

GATE	TENDER REQUIRES	FIRM HAS	RESULT
PEC category	≥ C-4	C-3	● Pass
Active filer	Yes	Yes	● Pass
Avg. turnover (3 yr)	≥ Rs 300M	Rs 410M	● Pass
Similar projects	3 × ≥ Rs 150M	2 on file	● Gap — add 1 completion cert
Bid security	Rs 12M BG	Capacity Rs 25M	● Pass
Non-blacklisting affidavit	Required	Not generated	● Gap — AI can draft now
ISO 9001	Mandatory	Not held	● Fail — cannot satisfy in time

Verdict line the user sees: **"Eligible with 2 fixable gaps — predicted technical score ≈ 74/100. Resolve both to bid safely. One hard requirement (ISO 9001) is unmet; review before proceeding."**

> **Closing note.** This blueprint is intentionally deep on the domain and the moat, because that — not the UI — is where this product is won. Build the data-capture flywheel first, keep the MVP to the wedge (eligibility → checklist → compiled PDF), and let agency-intelligence and the win/loss loop compound. That is how a $150/mo single-plan SaaS becomes something nobody can casually clone.